

ARX, RL, and LQR

Marco Conati

May 15, 2024

1 ARX

1.1 Introduction

In this task, we explore the identification of a dynamic system for a piezoelectric nozzle. Given time-domain step response data, we hope to model the system's behavior using ARX (AutoRegressive with eXogenous inputs) models. We are given external knowledge that the true model dynamics are between 5th and 10th order, we undertake an exhaustive search within this range to ascertain the best fitting model.

1.2 Methods

The ARX model is expressed as:

$$A(q)y(t) = B(q)u(t) + e(t)$$

where $y(t)$ is the system output, $u(t)$ is the system input, $e(t)$ is the noise term, and $A(q)$ and $B(q)$ are polynomials of the shift operator q^{-1} , representing the system's dynamics and how past inputs influence the current output. Specifically, $A(q) = 1 + a_1q^{-1} + \dots + a_{na}q^{-na}$ and $B(q) = b_1q^{-1} + \dots + b_{nb}q^{-nb}$, with na and nb indicating the orders of the respective polynomials.

The identification process involves estimating the coefficients of these polynomials such that the model output closely matches the observed output. This is achieved through the least squares problem, formulated as:

$$\hat{\theta} = \operatorname{argmin}_{\theta} \|Y - A * \theta\|^2$$

where $\theta = [a_1, \dots, a_{na}, b_1, \dots, b_{nb}]^T$ is the parameter vector, Y is the vector of observed outputs, and A is the regression matrix composed of past output and input values corresponding to the orders na and nb . This problem has solution:

$$\hat{\theta} = (A^T A)^{-1} A^T b$$

Where the first na elements of θ are the coefficients of a_{est} of $A(q)$, and the next nb elements are coefficients b_{est} of $B(q)$.

The matrices A and b are defined as:

$$A = \begin{bmatrix} -y(N_1 - 1) & \cdots & -y(N_1 - n) & u(N_1) & \cdots & u(N_1 - m) \\ \vdots & & \vdots & \vdots & & \vdots \\ -y(N_2 - 1) & \cdots & -y(N_2 - n) & u(N_2) & \cdots & u(N_2 - m) \\ \vdots & & \vdots & \vdots & & \vdots \end{bmatrix}$$

$$b = \begin{bmatrix} y(N_1) \\ \vdots \\ y(N_2) \\ \vdots \end{bmatrix}$$

with the amount of rows of A and b determined by the length of the data sample N . To identify the best-fitting model within the 5th to 10th order range, we first generate a random input sample of length N , and pass it through the unknown system to obtain the output y . Then, we iteratively construct ARX models for each combination of na and nb within the range $[5, 10]$, ensuring $na \geq nb$, and evaluate their performance based on the Normalized Root Mean Squared Error (NRMSE) between the predicted and actual outputs. The procedure is summarized in the following pseudocode:

```

1  for na = 5 to 10
2      for nb = 1 to na
3          Construct ARX model using orders na and nb
4          Estimate model parameters using least squares
5          Calculate NRMSE for the estimated model
6          If NRMSE < best_fit then
7              Update best_fit, best_na, and best_nb
8          End If
9      end for
10 end for
11 Report best_fit, best_na, and best_nb

```

Listing 1: Pseudocode for Model Identification

This systematic search allows us to pinpoint the ARX model that best encapsulates the dynamics of the `piezoNozzle` system as per the provided step-response data, `Vexample.mat`. In addition to this least-squares driven approach, we also estimate the dynamics using Matlab's `arx` function for comparison.

1.3 Results

The search for the optimal ARX model using my Least Squares implementation concluded with a 10th order model exhibiting a fit of 97.28% across 100 samples of random input and corresponding output data. The best fitting model coefficients for $A(z)$ and $B(z)$ are detailed as follows:

$$A(z) = 1 - 0.4224z^{-1} + 0.04663z^{-2} + 0.04079z^{-3} - 0.137z^{-4} - 0.04906z^{-5} \\ + 0.09359z^{-6} - 0.05533z^{-7} - 0.3532z^{-8} + 0.02261z^{-9} - 0.05159z^{-10}$$

$$B(z) = 0.01236z^{-1} - 0.09425z^{-2} - 0.1666z^{-3} + 0.2424z^{-4} + 0.1158z^{-5} \\ - 0.1463z^{-6} - 0.01962z^{-7} + 0.1386z^{-8} + 0.001422z^{-9} - 0.07413z^{-10}$$

However, the presence of higher-order terms with very small coefficients suggests potential overfitting. We can lower the order of the model to 8th order without a significant impact to the fit. This new model has a fit of 96.85%:

$$A(z) = 1 - 0.9407z^{-1} + 0.6646z^{-2} + 0.03094z^{-3} - 0.39z^{-4} + 0.5882z^{-5} \\ - 0.2344z^{-6} + 0.0227z^{-7} - 0.1296z^{-8}$$

$$B(z) = 0.01084z^{-1} - 0.1011z^{-2} - 0.1136z^{-3} + 0.2912z^{-4} - 0.09531z^{-5} \\ - 0.1213z^{-6} + 0.133z^{-7}$$

Parallely, employing MATLAB's System Identification Toolbox to construct an ARX model with the same order results in a ver similar model, albeit with a slightly better 97.35% fit:

$$A(z) = 1 - 0.9492z^{-1} + 0.6695z^{-2} + 0.0253z^{-3} - 0.3864z^{-4} + 0.5872z^{-5} \\ - 0.2349z^{-6} + 0.02091z^{-7} - 0.1315z^{-8}$$

$$B(z) = 0.002821 + 0.01021z^{-1} - 0.1011z^{-2} - 0.1134z^{-3} + 0.2932z^{-4} \\ - 0.09814z^{-5} - 0.121z^{-6} + 0.1327z^{-7}$$

The comparative performance of both models against the true system data is graphically represented in Figure 1, demonstrating their capability to closely track the actual data.




Figure 1: Performance of the Least Squares and MATLAB System Identification Toolbox ARX models against the true system data.

2 Problem 1b

2.1 Introduction and Assumptions

In this part of the study, the objective is to determine the appropriate voltage signal, $V^*(t)$, that should be input into a piezoelectric nozzle system to produce a specific desired pressure wave, P_{ref} , as dictated by data from the file $P_{\text{ref.mat}}$. Note that our ARX model from 1a) is linear, allowing the use of time reversal for approximating the adjoint operator G^* as $\tau G \tau e$. This is an important property for leveraging an Iterative Learning Control (ILC) approach. With this in mind, we can iteratively determine the signal $V^*(t)$ using our ARX model, and then compare the response of the true, hidden system $G(V^*(t))$ to the desired reference P_{ref} .

2.2 Methods

To derive the optimal voltage signal $V^*(t)$ for generating P_{ref} , I employed ILC with time reversal as follows:

1. Initialize V as a constant zero-voltage input and α as 0.1 (chosen heuristically). For iteration in maxIterations :

- (a) Compute the current pressure wave P_{current} by inputting V into the arx model for the system.
- (b) Calculate the error between P_{current} and P_{ref} .
- (c) Reverse the error signal in time using a time-reversal matrix τ .
- (d) Simulate the reversed error through the arx system to get P_{reversed} .
- (e) Apply a second time reversal to P_{reversed} to estimate the gradient of the error with respect to the voltage signal (adjoint effect). $G^*e = \tau G \tau e$
- (f) Update the voltage signal V according to $V_{i+1} = V_i + \alpha * G^*e$.

This process was repeated for 10,000 iterations, adjusting the voltage signal at each step to better match the desired pressure waveform. Note that G refers to the ARX model in this pseudocode.

2.3 Results

The iterative process resulted in a voltage signal $V^*(t)$ that successfully generated a pressure wave $G(V^*(t))$ closely matching P_{ref} within the 10,000 iterations. Importantly, the pressure wave $G(V^*(t))$ was generated by passing the signal $V^*(t)$ created using the ARX model through the hidden, real nozzle system. The reference pressure, achieved pressure, and corresponding input are shown in Figure 2.3.

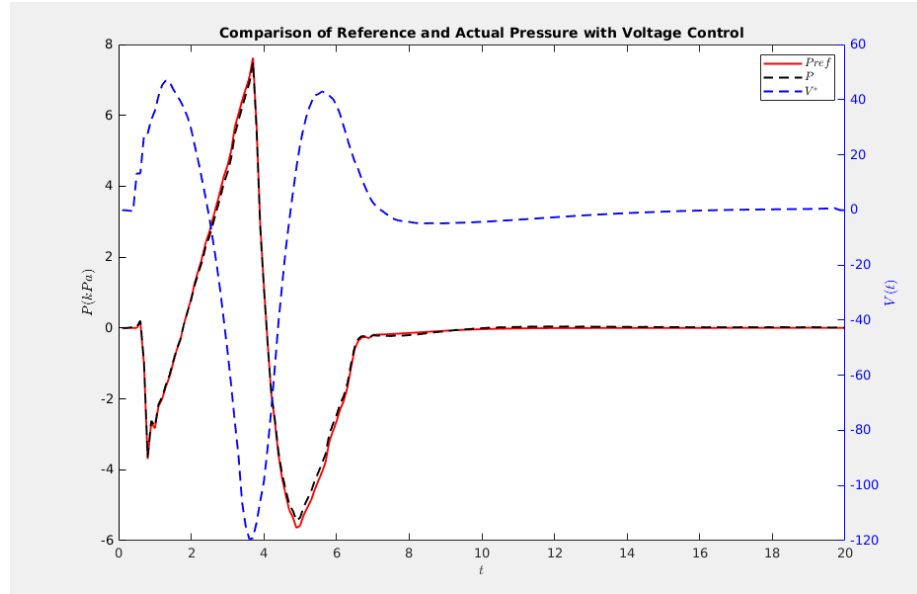


Figure 2: Reference pressure P_{ref} , achieved pressure $G(V^*(t))$, and corresponding input $V^*(t)$

3 Problem 2

3.1 Introduction

The wildlife park manager has tasked us with stabilizing the rabbit population at a constant level of 30 rabbits, despite the natural and predation-related fluctuations involving foxes. The populations are described by the Lotka-Volterra equations:

$$\begin{aligned}\frac{dR}{dt} &= \alpha R - \beta RF, \\ \frac{dF}{dt} &= \delta RF - \gamma F,\end{aligned}\tag{1}$$

where R and F denote the populations of rabbits and foxes, the challenge is to derive a control algorithm based on historical population data. Two methodologies are pursued: a Linear Quadratic Regulator (LQR) for managing the system within a linearized model, and Reinforcement Learning (RL) to iteratively discover an optimal control policy.

3.2 Methods - Parameter Identification

With both approaches, knowledge of the underlying population parameters is necessary. To determine population parameters for wildlife park's rabbit and fox populations, Autoregressive with exogenous inputs (ARX) models were fitted to historical data. The ARX model correlates the current output with past inputs and outputs, providing a framework suitable for systems influenced by prior states. In the context of the Lotka-Volterra predator-prey model, the interaction between rabbits and foxes suggests a specific structure for the ARX model's inputs.

Given the structure of the Lotka-Volterra equations:

$$\begin{aligned}\frac{dR}{dt} &= \alpha R - \beta RF, \\ \frac{dF}{dt} &= \delta RF - \gamma F,\end{aligned}$$

we can model each equation separately, treating the derivative terms as our state, and R , F , and RF terms as inputs. As an assumption, I approximated the population derivatives with a first-order difference. With this formulation, we get the following data for the rabbit model:

$$data_{rabbit} = iddata(dR_dt', [U_rabbit; U_rabbit * U_fox]', dt); \tag{2}$$

and for the fox population change model:

$$data_{fox} = iddata(dF_dt', [U_fox; U_rabbit * U_fox]', dt); \tag{3}$$

To fit the models, the following orders were set: $na = 0$ indicating no autoregressive component, $nb = [1, 1]$ reflecting the first-order dependence on both

input terms, and $nk = [0, 0]$ implying no delay. The ARX models are then computed using my Least Squares ARX implementation from task 1.

Model validation was performed by comparing the fitted models against the observed data, visually assessed by plotting the actual versus predicted population changes for both rabbits and foxes. The comparison plots are presented in Figures 1 and 2, demonstrating the models' fit to the historical data.

3.3 Results - Model Fitting and Validation

The ARX models fitted to the rabbit and fox populations yielded coefficients that align with the expected structure of the Lotka-Volterra equations. While the models do not perfectly capture the populations, they approximate the system's behavior generally. The initial parameters derived from the ARX models are $\alpha = 1.3$, $\beta = 0.089$, $\delta = 0.279$, and $\gamma = 0.009$, corresponding to the coefficients in the equations:

$$\begin{aligned}\frac{dR}{dt} &= \alpha R - \beta RF, \\ \frac{dF}{dt} &= \delta RF - \gamma F.\end{aligned}$$

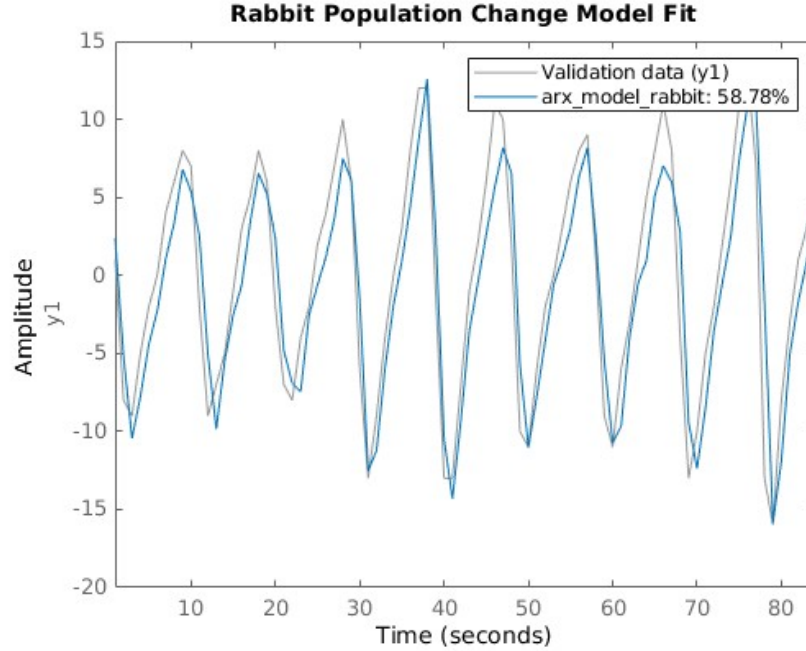


Figure 3: ARX model fit for rabbit population.

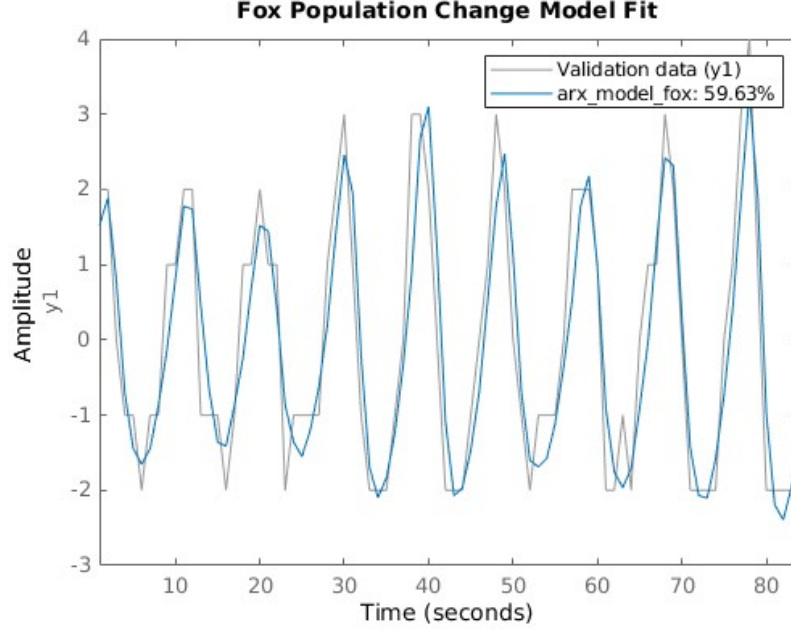


Figure 4: ARX model fit for fox population.

While the models are visually similar, the fits are not the most convincing. With this in mind, I tried refining these initial parameters using MATLAB's `fminsearch` function, which employs a derivative-free optimization algorithm based on the Nelder-Mead simplex method. This numerical optimization iteratively adjusts the parameters to minimize a cost function defined as the sum of squared errors between the observed populations and those predicted by integrating the Lotka-Volterra differential equations with the current parameter predictions:

$$\text{Cost Function} = \sum_{i=1}^n (R_{\text{observed},i} - R_{\text{predicted},i})^2 + (F_{\text{observed},i} - F_{\text{predicted},i})^2.$$

The refined model, with parameters $\alpha = 1.4787$, $\beta = 0.0990$, $\delta = 0.3025$, and $\gamma = 0.0103$, resulted in a lower cost function value, indicating a better fit to the data. The figures below showcase the improved fit of this refined model:

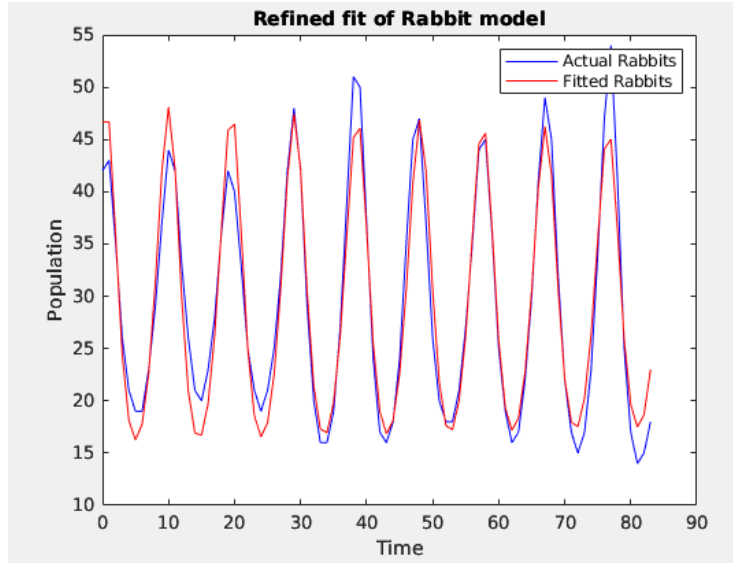


Figure 5: Refined model fit for rabbit population after numerical optimization.

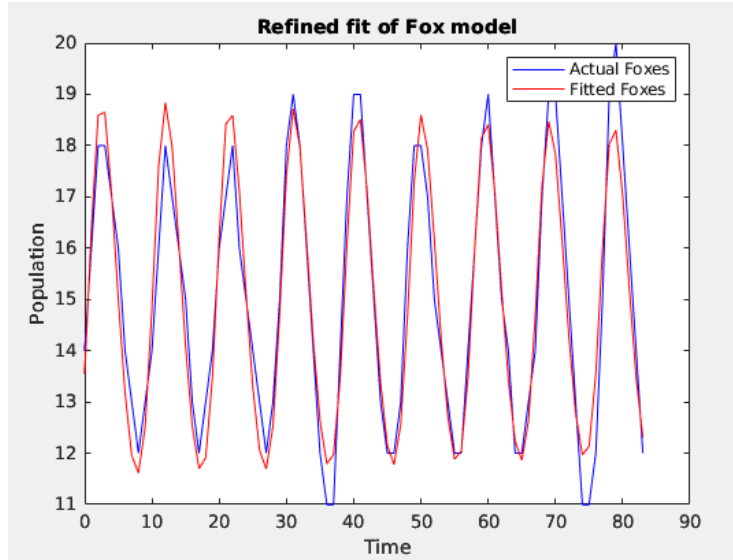


Figure 6: Refined model fit for fox population after numerical optimization.

Recall that the cost function, used by `fminsearch` for optimization, quantifies the aggregate discrepancy between the model's predictions and the actual data. The decrease in cost from 1011.53 to 711.785 confirms the better accuracy of the refined model, so it was used moving forward.

3.4 Controller 1 - Linearization and LQR

For the effective management of the rabbit and fox populations within the wildlife park, a control strategy employing Linear Quadratic Regulator (LQR) was developed. Two key assumptions were made with this algorithm:

1. All population values and changes were rounded up to whole numbers.
2. The nonlinear dynamics are closely approximated with a first-order linearization.

The linearization process involves calculating the Jacobian matrix, which represents a first-order approximation to the system's dynamics around a specified operating point. The Jacobian for the Lotka-Volterra system is given by:

$$J = \begin{bmatrix} \frac{\partial dR}{\partial R} & \frac{\partial dR}{\partial F} \\ \frac{\partial dF}{\partial R} & \frac{\partial dF}{\partial F} \end{bmatrix} \quad (4)$$

I set the desired state for the regulator at 30 rabbits and 15 foxes, as this is the target rabbit population with the closest whole number of foxes conducive to a balanced population.

Pseudocode for the LQR approach is provided below:

```
1 Set Lotka-Volterra parameters alpha, beta, delta, gamma
2 Define the equations dRdt and dFdt using Lotka-Volterra dynamics
3 Compute the Jacobian J of the system
4
5 Set the weight matrices Q for state and R_mat for control input
6 Q = [10 0; 0 1]; %Prioritize rabbit state
7 R_mat = 1;
8
9 for each simulation do
10     Initialize state with random whole numbers for R0 and F0
11     for each time step do
12         Compute the state matrix A using JacobianFunc at current
            state
13         Compute the LQR gain K using lqr(A, B, Q, R_mat)
14         Calculate the control input using the gain K and the state
            error
15         Apply control input, rounding up to whole numbers
16         Update the state using the Lotka-Volterra dynamics and dt
17         Round the updated state to whole numbers
18     end for
19 end for
```

Listing 2: Pseudocode for LQR Control Strategy

This pseudocode captures the essence of applying LQR control to a linearized system, with a focus on the continuous adjustment of the operating point and application of control inputs derived from LQR gains. Note that Q and R were set to emphasize the importance of maintaining the rabbit state, as this term has the most weight. With this formulation, the control inputs are designed to drive the state towards the desired population levels, while accounting for the discrete nature of the wildlife populations.

3.5 Controller 2 - RL

My project for MECHENG599 involves Double Deep Q-Networks (DDQN) for a different task, so I leveraged my implementation as an alternative approach to managing the wildlife park populations. DDQN extends the traditional Q-learning in two ways.

1. A traditional Q-Table is replaced with a NN for approximating values
2. Two separate NNs are used: a primary network for selecting the action and a target network for evaluation (this is the distinction between DQN and DDQN). This maintains more stability in the target, resulting in more stable training overall.

Mathematically, this is expressed with a new Bellman equation:

$$Q(s_t, a_t) = Q(s_t, a_t) + \alpha \left[r_{t+1} + \gamma \max_a Q'(s_{t+1}, a) - Q(s_t, a_t) \right], \quad (5)$$

where Q and Q' denote the primary and target networks, respectively, and γ is the discount factor.

My agent was designed as a 3-layer Multi-Layer Perceptron (MLP) with hidden layer sizes [64,128,64], which has access to the system state at every time step, facilitating the understanding and prediction of the population dynamics.

For the reinforcement learning to converge to an optimal policy, an appropriate reward function was crucial. The final reward function implemented was:

$$\begin{aligned} \text{rabbit_reward} &= -0.3 \times |\text{rabbits} - 30| \\ \text{fox_reward} &= -0.1 \times |\text{foxes} - 15| \\ \text{reward} &= 0.75 \times (\text{rabbit_reward} + \text{fox_reward}) \\ \text{if } |\text{rabbits} - 30| < 2 : \text{reward} &= 10 \end{aligned}$$

Termination conditions were set such that the simulation ends if the population of either species hits zero, exceeds the maximum allowed(1000), or if the maximum number of steps is reached. Moreover, a significant penalty was imposed for extinction or overpopulation events:

Penalty for Extinction or Overpopulation:

- If the rabbits' population is either zero or reaches the maximum population, or if the foxes' population reaches the maximum population:

$$\text{reward} = 1000 - (\text{current_step} \times 25)$$

This reward structure aims to avoid rapid extinction or overgrowth while densely encouraging the maintenance of the target populations. Coefficients were determined heuristically, with a few driving ideas:

1. The only positive reward should be achieved with a desirable outcome (In this case maintaining population within 2).

2. Maintaining rabbit population is more important than fox, in this case 3x as important
3. Surviving for one step without causing a termination condition should always be more beneficial than causing a termination condition (extinction or explosion).

The RL environment, along with the reward function, was set up using Python's `gym` and `torch` libraries, followed by the training of the agent to navigate the simulation towards the desired steady states of rabbit and fox populations.

3.6 Results - LQR and RL

The evaluation of the LQR controller involved initializing trajectories from randomized starting points within the ranges of [10-100] for both rabbits and foxes. The LQR method successfully guided the rabbit population towards the vicinity of the desired count of 30. While a small steady-state error persisted, the convergence of all populations remained within a margin of 2 rabbits, as depicted in Figures 7 and 8.

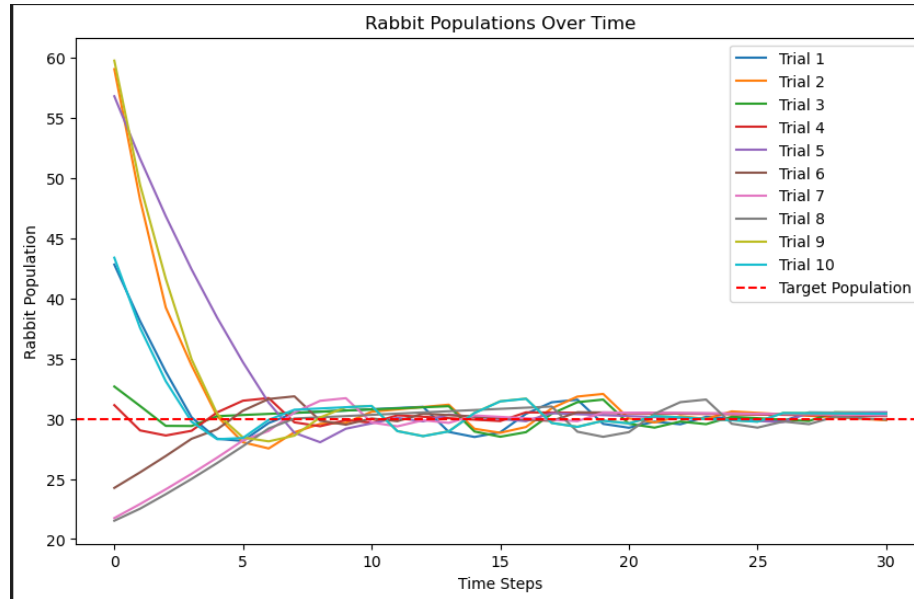


Figure 7: Rabbit population trajectories over time under LQR control.

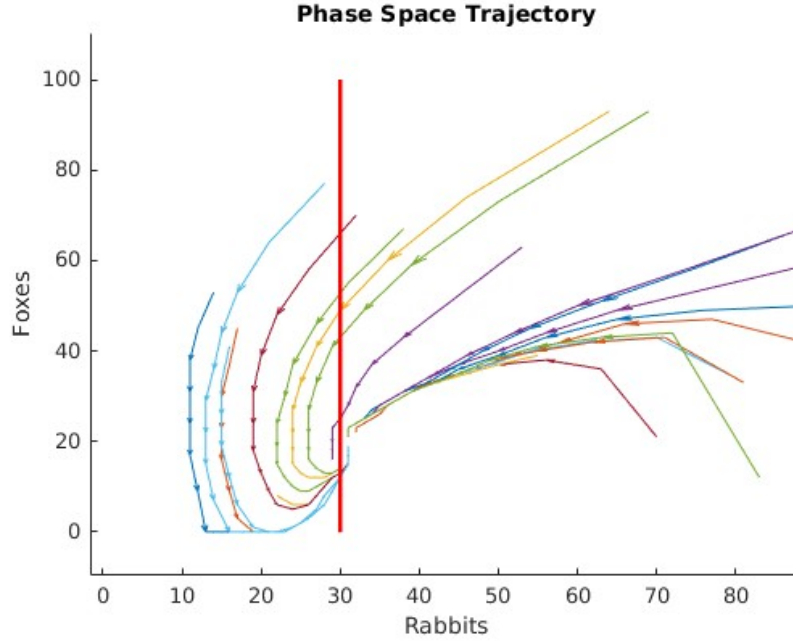


Figure 8: Population evolution in phase space with LQR control.

The RL-based controller also managed to achieve very similar convergence characteristics. Qualitatively, RL approach also appeared to be more stable at extreme starting locations, such as cases with 1000s of rabbits and few foxes, but numerically evaluating this goes beyond the scope of this task. The RL controller's performance with controlling the rabbit population is showcased in Figure 9, demonstrating its competent performance over a broad range of starting points.

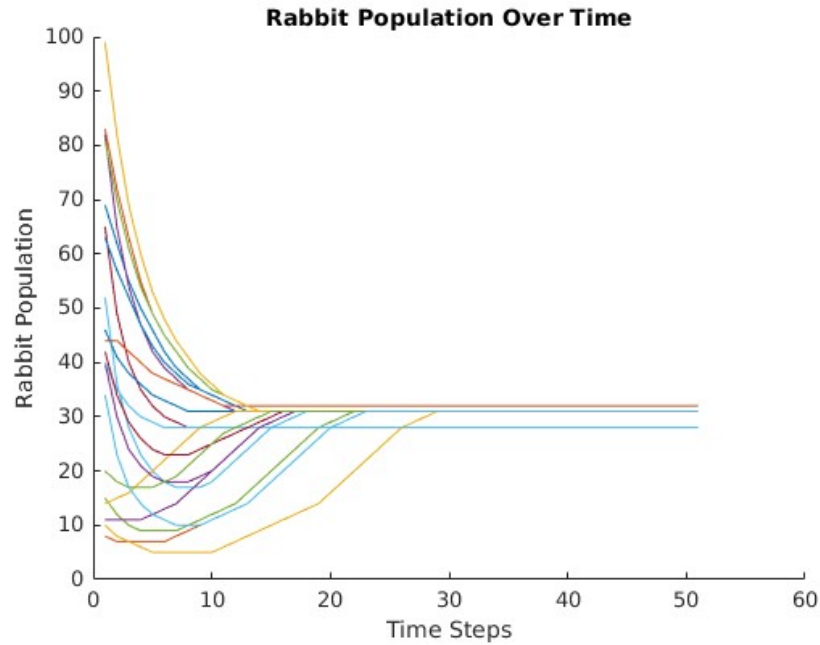


Figure 9: Rabbit population trajectories over time with RL control.

Overall, both control methodologies proved effective for randomized starting configurations within the [10-100] range for rabbits and foxes. With this in mind, both LQR and RL are viable solutions for addressing the wildlife park manager’s challenge.

Appendix

Python Code

```

1 import torch
2 import torch.nn as nn
3 import torch.optim as optim
4 import os
5 import numpy as np
6 import random
7 from collections import deque
8 from torch.nn.utils import clip_grad_norm_
9
10 # Check if CUDA is available and set the device accordingly
11 device = torch.device("cuda" if torch.cuda.is_available() else "cpu")
12
13 class QNetwork(nn.Module):
14     def __init__(self, input_dim, output_dim, hidden_layers):

```

```

15         super(QNetwork, self).__init__()
16         self.layers = nn.ModuleList()
17         for i in range(len(hidden_layers)):
18             self.layers.append(nn.Linear(input_dim if i == 0 else
hidden_layers[i-1], hidden_layers[i]))
19             self.layers.append(nn.ReLU())
20         self.layers.append(nn.Linear(hidden_layers[-1], output_dim)
)
21
22     def forward(self, x):
23         for layer in self.layers:
24             x = layer(x)
25         return x
26
27 class RLAgent:
28     def __init__(self, input_dim, action_dim, hidden_layers, lr=1e
-4, gamma=0.99, buffer_size=10000, batch_size=32,
29                 epsilon_start=1.0, epsilon_end=0.05, epsilon_decay
=0.995, target_update_frequency=10,
30                 learning_starts=1000, train_freq=4, max_grad_norm
=10):
31         #print(device)
32         self.q_network = QNetwork(input_dim, action_dim,
hidden_layers).to(device)
33         self.target_q_network = QNetwork(input_dim, action_dim,
hidden_layers).to(device)
34         self.target_q_network.load_state_dict(self.q_network.
state_dict())
35         self.optimizer = optim.Adam(self.q_network.parameters(), lr
=lr)
36         self.gamma = gamma
37         self.replay_buffer = deque(maxlen=buffer_size)
38         self.batch_size = batch_size
39         self.action_dim = action_dim
40         self.loss_fn = nn.MSELoss()
41         self.epsilon = epsilon_start
42         self.epsilon_end = epsilon_end
43         self.epsilon_decay = epsilon_decay
44         self.target_update_frequency = target_update_frequency
45         self.learning_starts = learning_starts
46         self.train_freq = train_freq
47         self.max_grad_norm = max_grad_norm
48
49     def update_replay_buffer(self, transition):
50         self.replay_buffer.append(transition)
51
52     def train(self):
53         if len(self.replay_buffer) < self.batch_size:
54             return
55         transitions = random.sample(self.replay_buffer, self.
batch_size)
56         state, action, next_state, reward, done = zip(*transitions)
57         state = torch.FloatTensor(np.array(state)).to(device)
58         next_state = torch.FloatTensor(np.array(next_state)).to(
device)
59         action = torch.LongTensor(action).squeeze().to(device)
60         reward = torch.FloatTensor(reward).to(device)

```

```

61     done = torch.FloatTensor(done).to(device)
62
63     current_q_values = self.q_network(state).gather(1, action.
64     unsqueeze(1)).squeeze(1)
65     ## Basic DQN
66     next_q_values = self.target_q_network(next_state).max(1)[0]
67     expected_q_values = reward + self.gamma * next_q_values *
68     (1 - done)
69     ##Double DQN
70     selected_actions = self.q_network(next_state).max(1)[1].
71     unsqueeze(1)
72     # Evaluate the selected action Q-value using the target
73     network
74     next_q_values = self.target_q_network(next_state).gather(1,
75     selected_actions).squeeze(1)
76     expected_q_values = reward + self.gamma * next_q_values.
77     detach() * (1 - done)
78
79     loss = self.loss_fn(current_q_values, expected_q_values.
80     detach())
81     self.optimizer.zero_grad()
82     loss.backward()
83
84     # Gradient clipping
85     clip_grad_norm_(self.q_network.parameters(), self.
86     max_grad_norm, norm_type=2)
87
88     self.optimizer.step()
89     return loss
90
91 def train_agent(self, env, num_episodes=1000, save_model_iters=
92 None, save_model_folder='/models', progress_bar = None, verbose
93 =False):
94     learning_starts = self.learning_starts
95     train_freq = self.train_freq
96     epsilon = self.epsilon
97     total_steps = 0
98     rewards = []
99     moving_averages = [] # Store moving averages
100
101     os.makedirs(save_model_folder, exist_ok=True)
102
103     for episode in range(num_episodes):
104         state = env.reset()
105         total_reward = 0
106         steps = 0
107         if progress_bar is not None:
108             progress_bar.update(1)
109         while True:
110             total_steps += 1
111             action = self.select_action(state, epsilon)
112             next_state, reward, done, _ = env.step(action)
113             total_reward += reward
114
115             self.update_replay_buffer((state, action,
116             next_state, reward, done))

```



```

107         if total_steps > learning_starts and total_steps %
train_freq == 0:
108             loss = self.train()
109
110             state = next_state
111             steps += 1
112
113             if done:
114                 break
115
116             rewards.append(total_reward)
117
118             # Calculate moving average of the last 100 rewards
119             if episode >= 99:
120                 moving_avg = sum(rewards[-100:]) / 100
121                 moving_averages.append(moving_avg)
122             else:
123                 # If less than 100 episodes, calculate the average
of what we have
124                 moving_avg = sum(rewards) / len(rewards)
125                 moving_averages.append(moving_avg)
126
127                 epsilon = max(self.epsilon_end, epsilon * self.
epsilon_decay)
128                 if episode % self.target_update_frequency == 0:
129                     self.update_target_network()
130
131                 if save_model_iters is not None and episode %
save_model_iters == 0:
132                     self.save_model(episode, save_model_folder, verbose
)
133                     print(f"Episode: {episode}, Total Reward: {
total_reward}, Steps: {steps}, Epsilon: {epsilon}, Moving Avg
Reward: {moving_avg}")
134                     if verbose:
135                         try:
136                             print(f"Loss: {loss.item()}, Episode: {episode
}, Total Reward: {total_reward}, Steps: {steps}, Epsilon: {
epsilon}, Moving Avg Reward: {moving_avg}")
137                         except:
138                             print(f"Episode: {episode}, Total Reward: {
total_reward}, Steps: {steps}, Epsilon: {epsilon}, Moving Avg
Reward: {moving_avg}")
139                     return moving_averages
140
141
142     def update_target_network(self):
143         self.target_q_network.load_state_dict(self.q_network.
state_dict())
144
145     def select_action(self, state, epsilon):
146         if random.random() > epsilon:
147             state = torch.FloatTensor(state).unsqueeze(0).to(device
) # Ensure state is on the correct device
148             q_value = self.q_network(state)
149             action = q_value.max(1)[1].item()
150         else:

```

```

151         action = random.randrange(self.action_dim)
152         return action
153
154     def save_model(self, episode, folder, verbose):
155         """Save the current model to disk within a specified folder
156         ."""
157         filename = os.path.join(folder, f'q_network_episode_{
158         episode}.pkl')
159         torch.save(self.q_network.state_dict(), filename)
160         if verbose:
161             print(f'Model saved to {filename}')
162
163     def load_model(self, episode, folder):
164         """Load a model from disk."""
165         filename = os.path.join(folder, f'q_network_episode_{
166         episode}.pkl')
167         if os.path.exists(filename):
168             self.q_network.load_state_dict(torch.load(filename,
169             map_location=device))
170             print(f'Model loaded from {filename}')
171         else:
172             print(f'No model file found at {filename}.')

```

Matlab Code

```

1 % FIT ARX MODEL TO 1A
2 N = 50;
3 dt = 1;
4 t = 0:dt:N-1;
5 rng(0);
6 u = 0.1 * randn(size(t));
7
8 % Obtain output from the piezo_nozzle model
9 P = piezo_nozzle(u, N); + 0.04
10
11 % Prepare the data
12 data = iddata(P, u', dt);
13
14 % Initialize variables to track the best models
15 best_fit_arx_ls = -Inf;
16 best_order_arx_ls = [];
17 best_a_est = [];
18 best_b_est = [];
19
20 best_fit_arx = -Inf;
21 best_order_arx = [];
22
23 % Test all 5th to 10th order models
24 for p = 8
25     for z = 7 % Ensure model has more poles than zeros or equal
26         % ARX Least Square Approach
27         [a_est, b_est] = arx_LS(u, P, t, p, p);
28         A = [1; a_est]; % Correct sign based on previous discussion
29         B = b_est;
30         C = 1;
31         model_arx_ls = idpoly(A', B', C);

```

```

32     y_pred_ls = predict(model_arx_ls, data);
33     fit_arx_ls = goodnessOfFit(y_pred_ls.y, data.y, 'NRMSE');
34
35     if fit_arx_ls > best_fit_arx_ls
36         best_fit_arx_ls = fit_arx_ls;
37         best_order_arx_ls_p = p;
38         best_order_arx_ls_z = z;
39         best_a_est = a_est;
40         best_b_est = b_est;
41         best_model_arx_ls = model_arx_ls;
42     end
43
44     % System Identification Toolbox Approach
45     model_order = [p z 1]; % nk = 1 to handle MATLAB indexing
46     (start from z^-1)
47     arx_model = arx(data, model_order);
48     y_pred_arx = predict(arx_model, data);
49     fit_arx = goodnessOfFit(y_pred_arx.y, data.y, 'NRMSE');
50
51     if fit_arx > best_fit_arx
52         best_fit_arx = fit_arx;
53         best_order_arx_p = p;
54         best_order_arx_z = z;
55         best_model_arx = arx_model;
56     end
57 end
58
59 % Report the best models
60 fprintf('Best ARX LS Model Order: %d, Fit: %.2f%%\n',
61         best_order_arx_ls, best_fit_arx_ls*100);
62 fprintf('Best ARX Model Order: %d, Fit: %.2f%%\n', best_order_arx,
63         best_fit_arx*100);
64
65 figure;
66
67 % Actual data
68 plot(data.y, 'k', 'LineWidth', 1.5);
69 hold on;
70
71 % Predictions from ARX model
72 y_pred_arx = predict(best_model_arx, data);
73 plot(y_pred_arx.y, 'r--', 'LineWidth', 1.2);
74
75 % Predictions from best ARX_LS model
76 y_pred_arx_ls = predict(best_model_arx_ls, data);
77 plot(y_pred_arx_ls.y, 'b-.', 'LineWidth', 1.2);
78
79 legend('Actual Data', 'ARX Model', 'ARX LS Model', 'Location', '
80 Best');
81 xlabel('Time');
82 ylabel('Output');
83 title('Comparison of Actual Data with ARX and ARX LS Models');
84 grid on;
85
86 function [a_est, b_est] = arx_LS(u, y, t, n, m)
87     N = length(y);

```

```

85     A = zeros(N-n, n+m);
86     b = zeros(N-n, 1);
87     for i = (n+1):N
88         y_segment = -y(i-1:-1:i-n)'; % Transpose to make it a row
            vector
89         u_segment = u(i:-1:i-m+1);
90
91         A(i-n, :) = [y_segment u_segment];
92         b(i-n) = y(i);
93     end
94     x = (A'*A) \ (A'*b);
95     a_est = x(1:n);
96     b_est = x(n+1:end);
97 end
98
99 function fit = goodnessOfFit(y_pred, y_true, method)
100     switch method
101     case 'NRMSE'
102         fit = 1 - norm(y_pred - y_true) / norm(y_true - mean(
            y_true));
103     otherwise
104         fit = NaN;
105     end
106 end
107 %Identify parameters of LotVol equations
108 load('fox_rabbit.mat')
109
110 u = (rabbit.*fox);
111 U = u(1:end-1);
112 U_rabbit = rabbit(1:end-1);
113 U_fox = fox(1:end-1);
114
115 data1 = iddata(dx_dt', [U_rabbit, U], dt);
116 data2 = iddata(dy_dt', [U_fox, U], dt);
117
118 na = 0;
119 nb = [1 1];
120 nk = [0 0];
121
122 model_order_rabbit = [na nb nk];
123 arx_model_rabbit = arx(data1, model_order_rabbit);
124
125 model_order_fox = [na nb nk];
126 arx_model_fox = arx(data2, model_order_fox);
127
128 figure;
129 compare(data1, arx_model_rabbit);
130 title('Rabbit Population Change Model Fit');
131
132 figure;
133 compare(data2, arx_model_fox);
134 title('Fox Population Change Model Fit');
135
136 load('fox_rabbit.mat')
137 td = [0:83];
138 rabbit = rabbit(1:85)';
139 fox = fox(1:85)';

```

```

140 p = [rabbit(1) fox(1) 1.5343    0.1025    0.2981    0.0101];
141
142 [p,fval,exitflag] = fminsearch(@leastcomp,p,[],td,rabbit,fox);
143
144
145 [t,y] = ode23(@lotvol,td,[p(1),p(2)],[],p(3),p(4),p(5),p(6));
146
147 figure;
148 plot(td, rabbit(1:end-1), 'b', td, y(:,1), 'r-');
149 xlabel('Time');
150 ylabel('Population');
151 legend('Actual Rabbits', 'Fitted Rabbits');
152 title('Refined fit of Rabbit model')
153
154 figure;
155 plot(td, fox(1:end-1), 'b', td, y(:,2), 'r-');
156 xlabel('Time');
157 ylabel('Population');
158 legend('Actual Foxes', 'Fitted Foxes');
159 title('Refined fit of Fox model')
160
161
162 function J = leastcomp(p,tdata,xdata,ydata)
163 n1 = length(tdata);
164 [t,y] = ode23(@lotvol,tdata,[p(1),p(2)],[],p(3),p(4),p(5),p(6));
165 errx = y(:,1)-xdata(1:n1)';
166 erry = y(:,2)-ydata(1:n1)';
167 J = errx'*errx + erry'*erry;
168 end
169
170 function dydt = lotvol(t,y,a1,a2,b1,b2)
171 tmp1 = a1*y(1) - a2*y(1)*y(2);
172 tmp2 = -b1*y(2) + b2*y(1)*y(2);
173 dydt = [tmp1; tmp2];
174 end
175
176
177 %Simulate LQR 2
178 syms R F real
179 alpha = 1.5343;
180 beta = 0.1025;
181 delta = 0.2981;
182 gamma = 0.0101;
183
184 dRdt = alpha * R - beta * R * F;
185 dFdt = gamma * R * F - delta * F;
186
187 J = jacobian([dRdt, dFdt], [R, F]);
188
189 JacobianFunc = matlabFunction(J, 'Vars', [R, F]);
190
191 Q = [0.1 0; 0 0.1];
192 R_mat = 10;
193
194 dt = 1;
195
196 colors = lines(20);

```

```

197
198 R_values_all = cell(1, 20);
199 F_values_all = cell(1, 20);
200
201 figure;
202 hold on;
203
204 for j = 1:20
205     R0 = ceil(rand * 10);
206     F0 = ceil(rand * 10);
207
208     state = [R0; F0];
209
210     R_values = R0;
211     F_values = F0;
212
213     for i = 1:50
214         disp(state)
215         A = double(JacobianFunc(state(1), state(2)));
216         B = [0; 1];
217
218         K = lqr(A, B, Q, R_mat);
219         control_input = -K * (state - [30; 15]);
220
221         state = state + B * ceil(control_input * dt);
222         state = ceil(state);
223         state = max(state, [0; 0]);
224
225         dPop = double([alpha * state(1) - beta * state(1) * state
226         (2); -delta * state(2) + gamma * state(1) * state(2)]);
227         state = state + dPop * dt;
228         state = ceil(state);
229         state = max(state, [0; 0]);
230
231         R_values = [R_values state(1)];
232         F_values = [F_values state(2)];
233     end
234
235     R_values_all{j} = R_values;
236     F_values_all{j} = F_values;
237
238     plot(R_values, F_values, 'Color', colors(j,:));
239     if j == 1 % Only label once to avoid duplicate labels in legend
240         quiver(R_values(1:end-1), F_values(1:end-1), diff(R_values)
241         , diff(F_values), 0, 'MaxHeadSize', 0.5, 'Color', colors(j,:),
242         'AutoScale', 'off');
243     end
244 end
245
246 xlabel('Rabbits');
247 ylabel('Foxes');
248 title('Phase Space Trajectory');
249 hold off;
250
251 % New figure for Rabbit population vs. Time
252 figure;
253 hold on;

```

```

251
252 for j = 1:20
253     plot(1:length(R_values_all{j}), R_values_all{j}, 'Color',
254         colors(j,:)); % Plot each trial's rabbit population over time
255 end
256 xlabel('Time Steps');
257 ylabel('Rabbit Population');
258 title('Rabbit Population Over Time');
259 ylim([0, 100]); % Fix the y-axis to be 0-100
260
261 hold off;

```